The model suggests, plain code validates, a person decides, and only then does anything happen.

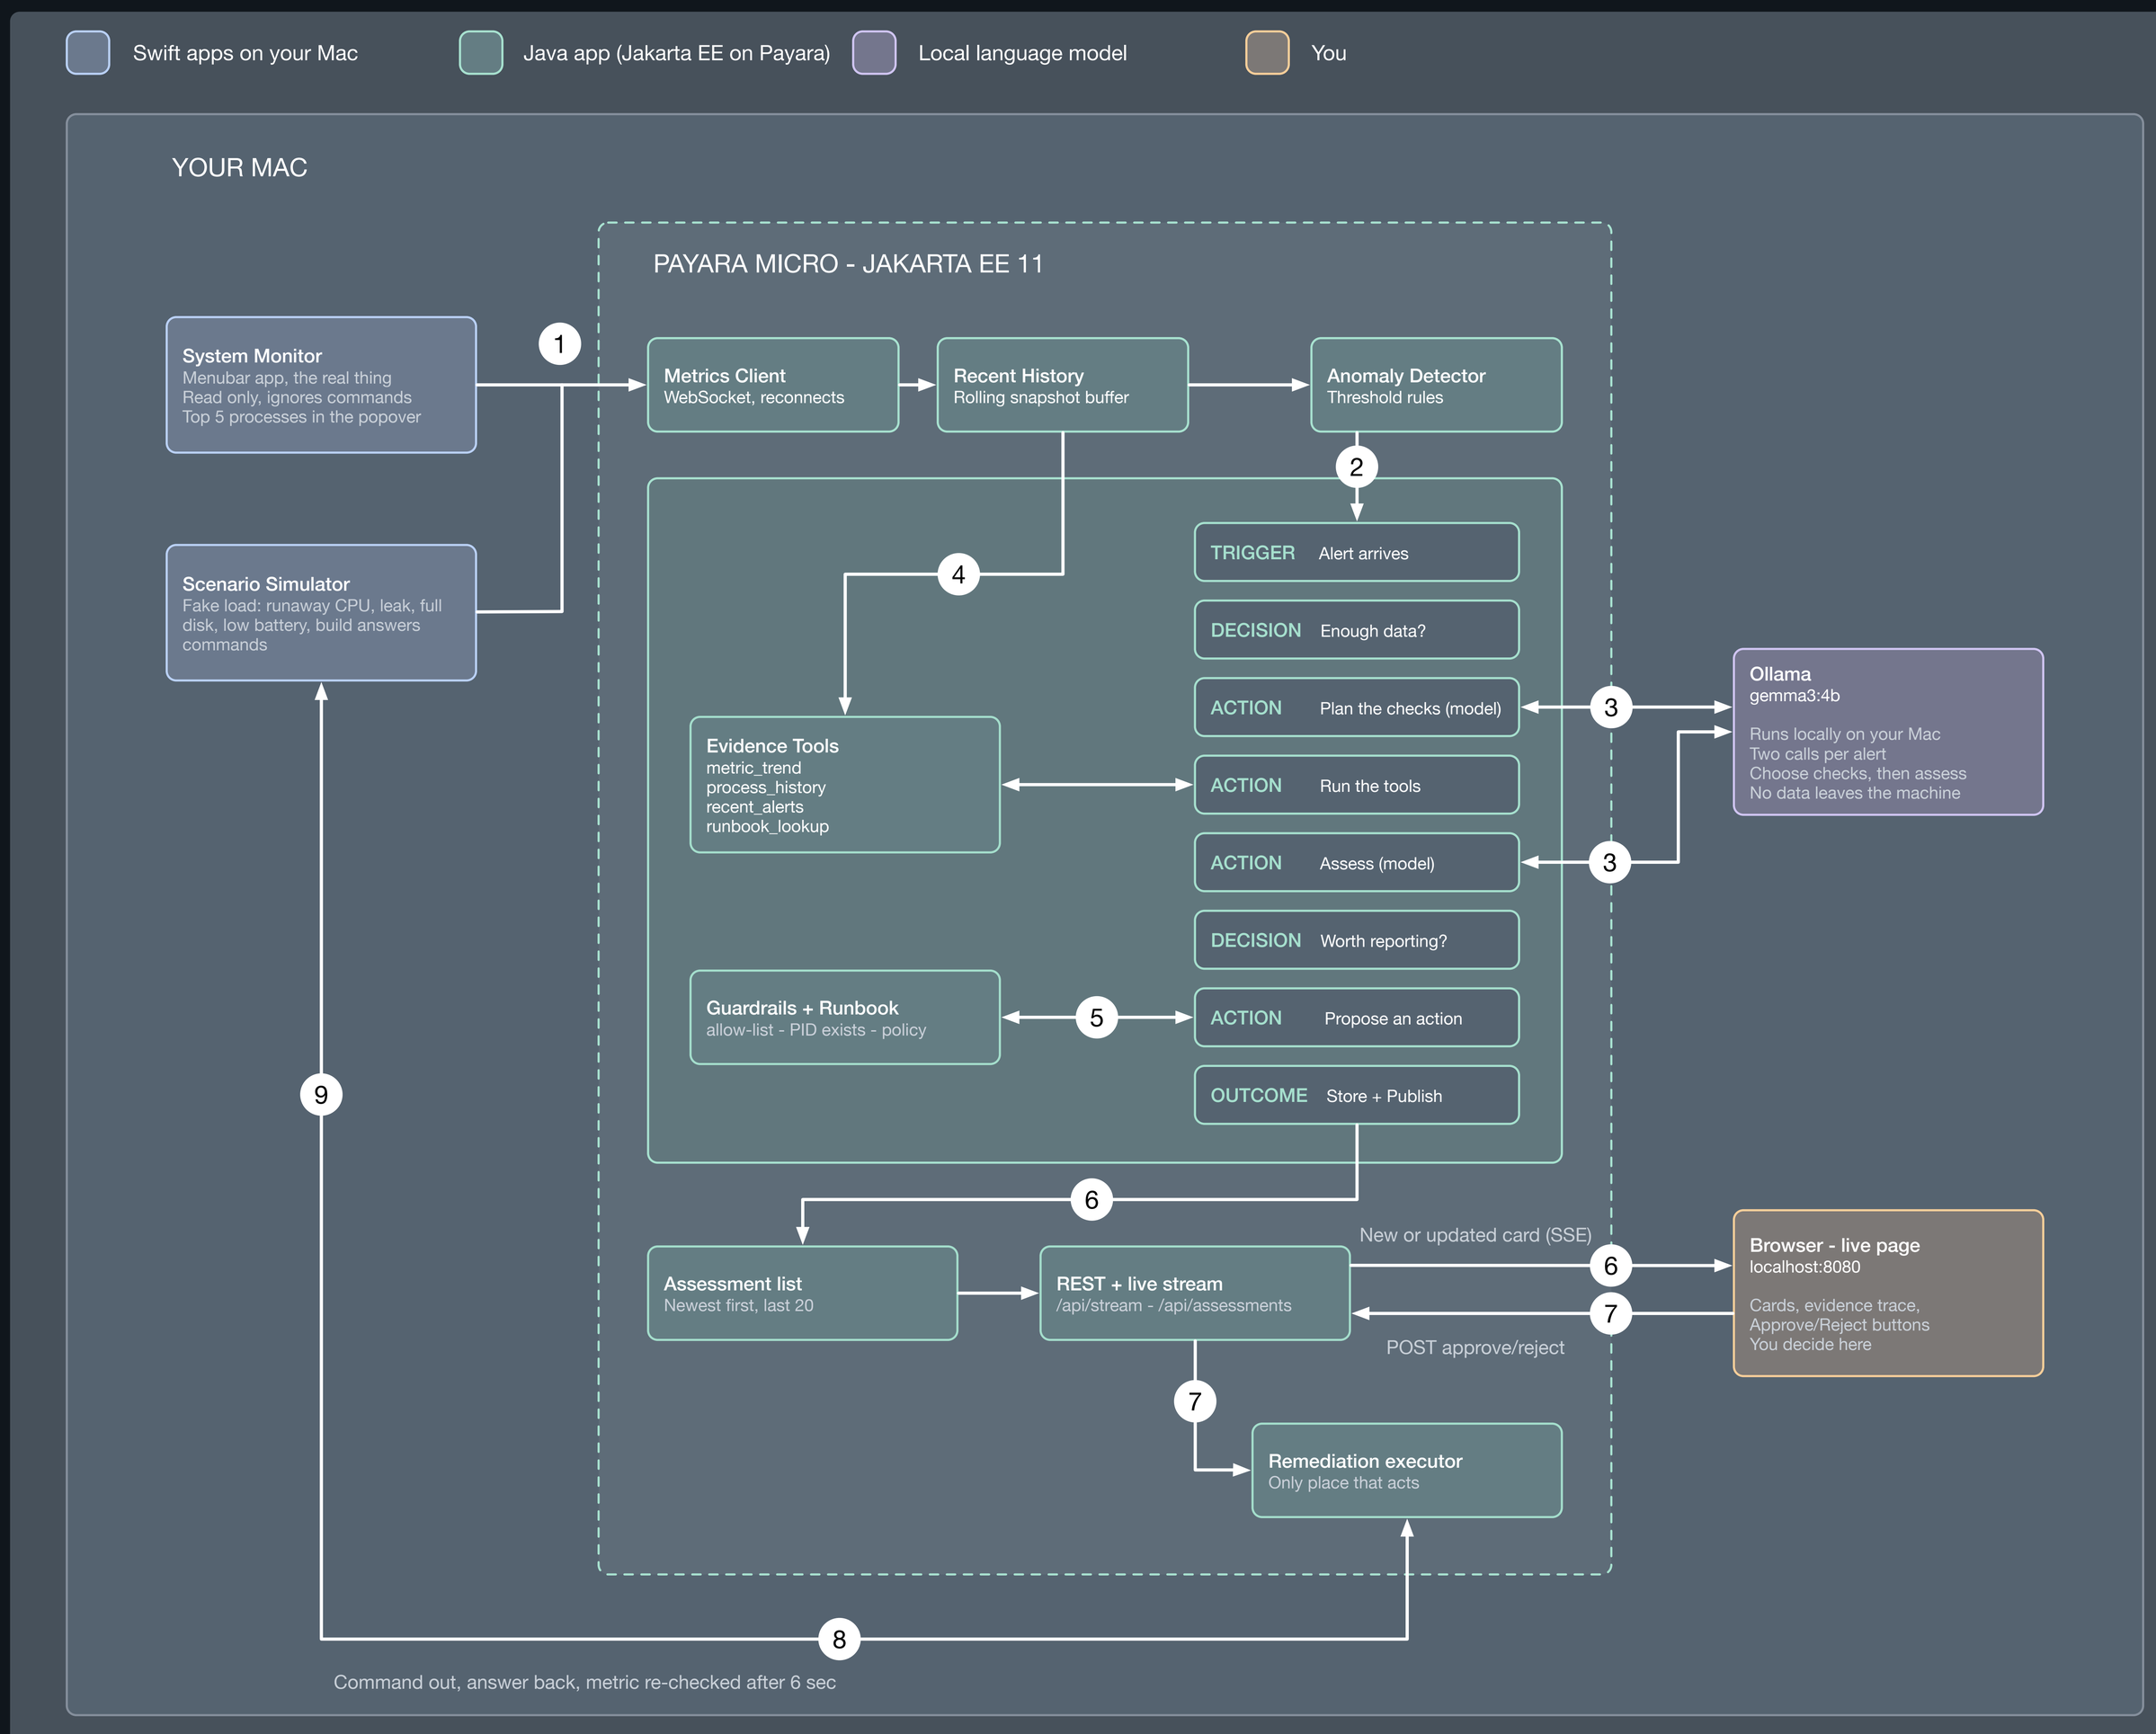

1 **A snapshot arrives every second**

The System Monitor (or the Scenario Simulator) sends CPU, memory, network, disk, battery and the top processes as JSON over a local WebSocket

2 **The detector fires an alert**

When a threshold holds across the time window, an alert event starts the agent. A cooldown stops repeat alerts.

3 **The model plans, then assesses**

First call: which evidence tools are worth running. Second call: severity, likely cause, recommendation and an optional action, based on what the tools found.

4 **Tools gather evidence**

Trend of the metric, process history, earlier alerts and the team's runbook. Plain code reads the recent history; the model never touches the system.

5 **Guardrails check the proposal**

The action must be on the allow-list, the process must exist, and the runbook must permit it. Protected system processes and build tools are blocked.

6 **The result is stored and pushed**

The assessment goes into the list and appears on the live page with its evidence trace and, if allowed, an Approve button.

7 **You approve or reject**

One click claims the proposal, so a double click cannot run it twice. Approve hands it to the remediation executor.

8 **The command goes to the simulator**

Terminate a process or purge temporary files, sent over the same WebSocket the data came in on.

9 **Answer, then verification**

The simulator confirms. Six seconds later the executor measures again and the card shows whether the metric actually dropped.